

```
In [1]: import numpy as np # for linear algebra
import pandas as pd

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

```
In [2]: import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')
```

```
In [3]: df=pd.read_csv('creditcard.csv')
df.head(10)
```

```
Out[3]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.0
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.0
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.0
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.0
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.0
5	2.0	-0.425966	0.960523	1.141109	-0.168252	0.420987	-0.029728	0.476201	0.260314	-0.0
6	4.0	1.229658	0.141004	0.045371	1.202613	0.191881	0.272708	-0.005159	0.081213	0.0
7	7.0	-0.644269	1.417964	1.074380	-0.492199	0.948934	0.428118	1.120631	-3.807864	0.0
8	7.0	-0.894286	0.286157	-0.113192	-0.271526	2.669599	3.721818	0.370145	0.851084	-0.0
9	9.0	-0.338262	1.119593	1.044367	-0.222187	0.499361	-0.246761	0.651583	0.069539	-0.0

10 rows × 31 columns

```
In [4]: df.tail(10)
```

Out[4]:

	Time	V1	V2	V3	V4	V5	V6	V7	
284797	172782.0	-0.241923	0.712247	0.399806	-0.463406	0.244531	-1.343668	0.929369	-0.2
284798	172782.0	0.219529	0.881246	-0.635891	0.960928	-0.152971	-1.014307	0.427126	0.1
284799	172783.0	-1.775135	-0.004235	1.189786	0.331096	1.196063	5.519980	-1.518185	2.0
284800	172784.0	2.039560	-0.175233	-1.196825	0.234580	-0.008713	-0.726571	0.017050	-0.1
284801	172785.0	0.120316	0.931005	-0.546012	-0.745097	1.130314	-0.235973	0.812722	0.1
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.918215	7.3
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.024330	0.2
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.296827	0.7
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.686180	0.6
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.577006	-0.4

10 rows × 31 columns

In [5]: `df.shape`

Out[5]: (284807, 31)

In [6]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype  
---  -
0   Time        284807 non-null float64
1   V1          284807 non-null float64
2   V2          284807 non-null float64
3   V3          284807 non-null float64
4   V4          284807 non-null float64
5   V5          284807 non-null float64
6   V6          284807 non-null float64
7   V7          284807 non-null float64
8   V8          284807 non-null float64
9   V9          284807 non-null float64
10  V10         284807 non-null float64
11  V11         284807 non-null float64
12  V12         284807 non-null float64
13  V13         284807 non-null float64
14  V14         284807 non-null float64
15  V15         284807 non-null float64
16  V16         284807 non-null float64
17  V17         284807 non-null float64
18  V18         284807 non-null float64
19  V19         284807 non-null float64
20  V20         284807 non-null float64
21  V21         284807 non-null float64
22  V22         284807 non-null float64
23  V23         284807 non-null float64
24  V24         284807 non-null float64
25  V25         284807 non-null float64
26  V26         284807 non-null float64
27  V27         284807 non-null float64
28  V28         284807 non-null float64
29  Amount      284807 non-null float64
30  Class       284807 non-null int64  
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

```
In [7]: df["Class"].value_counts()
```

```
Out[7]: 0      284315
        1        492
        Name: Class, dtype: int64
```

```
In [8]: df = df.drop(['Time'],axis=1)
        df.head()
```

```
Out[8]:
```

	V1	V2	V3	V4	V5	V6	V7	V8	V9
0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787
1	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425
2	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654
3	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024
4	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739

5 rows × 30 columns

```
In [9]: df.shape
```

```
Out[9]: (284807, 30)
```


```
In [10]: df.duplicated().any()
```

```
Out[10]: True
```

```
In [11]: df = df.drop_duplicates()
df.shape
```


```
Out[11]: (275663, 30)
```

```
In [12]: df["Class"].value_counts(normalize=True).plot(
        kind="bar", xlabel="Class", ylabel="Relative Frequency", title="Class Balance"
    );
```



In [13]: `df.hist(bins=30, figsize=(30, 30))`

Out[13]: `array([[<Axes: title={'center': 'V1'}>, <Axes: title={'center': 'V2'}>, <Axes: title={'center': 'V3'}>, <Axes: title={'center': 'V4'}>, <Axes: title={'center': 'V5'}>], [<Axes: title={'center': 'V6'}>, <Axes: title={'center': 'V7'}>, <Axes: title={'center': 'V8'}>, <Axes: title={'center': 'V9'}>, <Axes: title={'center': 'V10'}>], [<Axes: title={'center': 'V11'}>, <Axes: title={'center': 'V12'}>, <Axes: title={'center': 'V13'}>, <Axes: title={'center': 'V14'}>, <Axes: title={'center': 'V15'}>], [<Axes: title={'center': 'V16'}>, <Axes: title={'center': 'V17'}>, <Axes: title={'center': 'V18'}>, <Axes: title={'center': 'V19'}>, <Axes: title={'center': 'V20'}>], [<Axes: title={'center': 'V21'}>, <Axes: title={'center': 'V22'}>, <Axes: title={'center': 'V23'}>, <Axes: title={'center': 'V24'}>, <Axes: title={'center': 'V25'}>], [<Axes: title={'center': 'V26'}>, <Axes: title={'center': 'V27'}>, <Axes: title={'center': 'V28'}>, <Axes: title={'center': 'Amount'}>, <Axes: title={'center': 'Class'}>]], dtype=object)`




```
In [14]: df.describe()
```

Out[14]:	V1	V2	V3	V4	V5	V6
count	275663.000000	275663.000000	275663.000000	275663.000000	275663.000000	275663.000000
mean	-0.037460	-0.002430	0.025520	-0.004359	-0.010660	-0.014206
std	1.952522	1.667260	1.507538	1.424323	1.378117	1.313213
min	-56.407510	-72.715728	-48.325589	-5.683171	-113.743307	-26.160506
25%	-0.941105	-0.614040	-0.843168	-0.862847	-0.700192	-0.765861
50%	-0.059659	0.070249	0.200736	-0.035098	-0.060556	-0.270931
75%	1.294471	0.819067	1.048461	0.753943	0.604521	0.387704
max	2.454930	22.057729	9.382558	16.875344	34.801666	73.301626

8 rows × 30 columns

Scaling the dataset

```
In [15]: from sklearn.preprocessing import RobustScaler
new_df = df.copy()
new_df['Amount'] = RobustScaler().fit_transform(new_df['Amount'].to_numpy().reshape(
new_df['Amount'].hist();
```



```
In [16]: new_df['Amount'].describe()
```

```
Out[16]: count    275663.000000
mean         0.908007
std          3.439940
min         -0.322511
25%         -0.236924
50%          0.000000
75%          0.763076
max          348.694743
Name: Amount, dtype: float64
```

Copying the contents of the data into new_df

```
In [17]: new_df.head()
```

```
Out[17]:
```

	V1	V2	V3	V4	V5	V6	V7	V8	V9
0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787
1	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425
2	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654
3	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024
4	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739

5 rows × 30 columns

Splitting the dataset into training and testing data

```
In [18]: X = new_df.drop('Class',axis=1)
         y = new_df['Class']
```

```
In [19]: from sklearn.model_selection import train_test_split

X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.20, random_state=4)

print("X_train shape:", X_train.shape)
print("y_train shape:", y_train.shape)
print("X_test shape:", X_test.shape)
print("y_test shape:", y_test.shape)
```

X_train shape: (220530, 29)

y_train shape: (220530,)

X_test shape: (55133, 29)

y_test shape: (55133,)

Building Models with the unbalanced dataset

Logistic Regression

```
In [20]: from sklearn.linear_model import LogisticRegression

logistic_model = LogisticRegression(max_iter=1000)
logistic_model.fit(X_train, y_train)
logistic_model.score(X_train, y_train)
```

```
Out[20]: 0.999183784519113
```

```
In [21]: # Predicting the result
         y_pred = logistic_model.predict(X_test)
```

```
In [22]: from sklearn.metrics import accuracy_score,confusion_matrix,classification_report

ac = accuracy_score(y_test,y_pred)*100
cm = confusion_matrix(y_test,y_pred)
cr= classification_report(y_test,y_pred, target_names=['Not Fraud', 'Fraud'])
print("accuracy score:",ac)
print("confusion matrix:",cm)
print("classification report:",cr)
```



```

accuracy score: 99.92200678359603
confusion matrix: [[55035    7]
 [   36   55]]
classification report:

```

			precision	recall	f1-score	support
Not Fraud	1.00	1.00	1.00	55042		
Fraud	0.89	0.60	0.72	91		
accuracy			1.00	55133		
macro avg	0.94	0.80	0.86	55133		
weighted avg	1.00	1.00	1.00	55133		

Random Forest

```

In [23]: from sklearn.ensemble import RandomForestClassifier

rf = RandomForestClassifier(max_depth=2, n_jobs=-1)
rf.fit(X_train, y_train)

```

```

Out[23]: ▼          RandomForestClassifier
RandomForestClassifier(max_depth=2, n_jobs=-1)

```

```

In [24]: y_pred=rf.predict(X_test)

```

```

In [25]: ac = accuracy_score(y_test,y_pred)*100
cm = confusion_matrix(y_test,y_pred)
cr= classification_report(y_test,y_pred, target_names=['Not Fraud', 'Fraud'])
print("accuracy score:",ac)
print("confusion matrix:",cm)
print("classification report:",cr)

```

```

accuracy score: 99.89842743910181
confusion matrix: [[55030    12]
 [   44   47]]
classification report:

```

			precision	recall	f1-score	support
Not Fraud	1.00	1.00	1.00	55042		
Fraud	0.80	0.52	0.63	91		
accuracy			1.00	55133		
macro avg	0.90	0.76	0.81	55133		
weighted avg	1.00	1.00	1.00	55133		

Naive Bayes GaussianNB

In [26]: `from sklearn.naive_bayes import GaussianNB`

```
gnb = GaussianNB()
gnb.fit(X_train, y_train)

y_pred = gnb.predict(X_test)

ac = accuracy_score(y_test,y_pred)*100
cm = confusion_matrix(y_test,y_pred)
cr= classification_report(y_test,y_pred, target_names=['Not Fraud', 'Fraud'])
print("accuracy score:",ac)
print("confusion matrix:",cm)
print("classification report:",cr)
```

accuracy score: 97.81618994068889

confusion matrix: [[53857 1185]

[19 72]]

classification report: precision recall f1-score support

Not Fraud	1.00	0.98	0.99	55042
-----------	------	------	------	-------

Fraud	0.06	0.79	0.11	91
-------	------	------	------	----

accuracy			0.98	55133
----------	--	--	------	-------

macro avg	0.53	0.88	0.55	55133
-----------	------	------	------	-------

weighted avg	1.00	0.98	0.99	55133
--------------	------	------	------	-------

Decision Tree

In [27]: `from sklearn.tree import DecisionTreeClassifier`

```
dtc = DecisionTreeClassifier(random_state=42)
dtc.fit(X_train, y_train)

y_pred = dtc.predict(X_test)

ac = accuracy_score(y_test,y_pred)*100
cm = confusion_matrix(y_test,y_pred)
cr= classification_report(y_test,y_pred, target_names=['Not Fraud', 'Fraud'])
print("accuracy score:",ac)
print("confusion matrix:",cm)
print("classification report:",cr)
```

accuracy score: 99.89479984764115

confusion matrix: [[55007 35]

[23 68]]

classification report: precision recall f1-score support

Not Fraud	1.00	1.00	1.00	55042
-----------	------	------	------	-------

Fraud	0.66	0.75	0.70	91
-------	------	------	------	----

accuracy			1.00	55133
----------	--	--	------	-------

macro avg	0.83	0.87	0.85	55133
-----------	------	------	------	-------

weighted avg	1.00	1.00	1.00	55133
--------------	------	------	------	-------

Balanced Dataset with OverSampling Technique

```
In [28]: X = new_df.drop('Class',axis=1)
y = new_df['Class']

print("X.shape: ", X.shape)
print("y.shape: ", y.shape)
```

```
X.shape: (275663, 29)
y.shape: (275663,)
```

```
In [29]: from imblearn.over_sampling import SMOTE

X_res,y_res = SMOTE().fit_resample(X,y)
y_res.value_counts()
```

```
Out[29]: 0    275190
1    275190
Name: Class, dtype: int64
```

Train Test Split on Balanced data

```
In [30]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X_res,y_res,test_size=0.20,random_

print("X_train shape:", X_train.shape)
print("y_train shape:", y_train.shape)
print("X_test shape:", X_test.shape)
print("y_test shape:", y_test.shape)
```

```
X_train shape: (440304, 29)
y_train shape: (440304,)
X_test shape: (110076, 29)
y_test shape: (110076,)
```

Logistic Regression on Balanced Data

```
In [31]: from sklearn.linear_model import LogisticRegression

logistic_model = LogisticRegression(max_iter=1000)
logistic_model.fit(X_train, y_train)

y_pred = logistic_model.predict(X_test)

ac = accuracy_score(y_test,y_pred)*100
cm = confusion_matrix(y_test,y_pred)
cr= classification_report(y_test,y_pred, target_names=['Not Fraud', 'Fraud'])
print("accuracy score:",ac)
print("confusion matrix:",cm)
print("classification report:",cr)
```

```

accuracy score: 94.49925506014026
confusion matrix: [[53731  1342]
 [ 4713 50290]]
classification report:

```

			precision	recall	f1-score	support
Not Fraud	0.92	0.98	0.95	55073		
Fraud	0.97	0.91	0.94	55003		
accuracy			0.94	110076		
macro avg	0.95	0.94	0.94	110076		
weighted avg	0.95	0.94	0.94	110076		

Random Forest on Balanced Data

```

In [32]: from sklearn.ensemble import RandomForestClassifier

rf = RandomForestClassifier(max_depth=2, n_jobs=-1)
rf.fit(X_train, y_train)

y_pred = rf.predict(X_test)

ac = accuracy_score(y_test,y_pred)*100
cm = confusion_matrix(y_test,y_pred)
cr= classification_report(y_test,y_pred, target_names=['Not Fraud', 'Fraud'])
print("accuracy score:",ac)
print("confusion matrix:",cm)
print("classification report:",cr)

```

```

accuracy score: 92.85947890548348
confusion matrix: [[54708  365]
 [ 7495 47508]]
classification report:

```

			precision	recall	f1-score	support
Not Fraud	0.88	0.99	0.93	55073		
Fraud	0.99	0.86	0.92	55003		
accuracy			0.93	110076		
macro avg	0.94	0.93	0.93	110076		
weighted avg	0.94	0.93	0.93	110076		

GaussianNB on Balanced data

```

In [33]: from sklearn.naive_bayes import GaussianNB

gnb = GaussianNB()
gnb.fit(X_train, y_train)

y_pred = gnb.predict(X_test)

ac = accuracy_score(y_test,y_pred)*100
cm = confusion_matrix(y_test,y_pred)
cr= classification_report(y_test,y_pred, target_names=['Not Fraud', 'Fraud'])
print("accuracy score:",ac)
print("confusion matrix:",cm)
print("classification report:",cr)

```

```
accuracy score: 91.18336422108362
confusion matrix: [[53698  1375]
 [ 8330 46673]]
classification report:
```

			precision	recall	f1-score	support
Not Fraud	0.87	0.98	0.92	55073		
Fraud	0.97	0.85	0.91	55003		
accuracy			0.91	110076		
macro avg	0.92	0.91	0.91	110076		
weighted avg	0.92	0.91	0.91	110076		

Decision Tree on Balanced data

```
In [34]: from sklearn.tree import DecisionTreeClassifier

dtc = DecisionTreeClassifier(max_depth=6, random_state=42)
dtc.fit(X_train, y_train)

y_pred = dtc.predict(X_test)

ac = accuracy_score(y_test,y_pred)*100
cm = confusion_matrix(y_test,y_pred)
cr= classification_report(y_test,y_pred, target_names=['Not Fraud', 'Fraud'])
print("accuracy score:",ac)
print("confusion matrix:",cm)
print("classification report:",cr)
```

```
accuracy score: 96.14629892074566
confusion matrix: [[52862  2211]
 [ 2031 52972]]
classification report:
```

			precision	recall	f1-score	support
Not Fraud	0.96	0.96	0.96	55073		
Fraud	0.96	0.96	0.96	55003		
accuracy			0.96	110076		
macro avg	0.96	0.96	0.96	110076		
weighted avg	0.96	0.96	0.96	110076		

In []: